

What Causes Binge Drinking?

Statistical Techniques for Identifying At-Risk Students

Yuliya Dovzhenko, Brandon Podmayersky, James Yang
Princeton University

Abstract

Text of the abstract.

1 Introduction

Alcohol is a depressant drug that, when used in moderation, acts as a relaxant and social lubricant. Responsible use of alcohol for recreation is a fundamental part of American culture and many other cultures worldwide, but heavy use can have serious consequences. Alcohol abuse is a leading factor in violent crime and automotive fatalities, alcohol poisoning can be fatal in a short time period, and chronic alcoholism can lead to long-term liver failure and reduced brain function.

Binge drinking and overconsumption of alcohol are particularly widespread problems on college campuses, where many students are exposed to serious health risks from drinking. Efforts targeting the elimination of alcohol consumption completely have historically failed and represent a grievous threat to individual liberties. Instead, trends have pointed to education and healthy drinking practices to reduce the danger to students.

One thing that seems clear, though, is that not all individuals are at an equal risk of problem drinking. If we could find some way to identify students that are most likely to engage in risky drinking behavior, alcohol education programs and other initiatives could be targeted specifically at these groups. In this work we use survey data from the 2001 Harvard School of Public Health College Alcohol Study to show that information about a student's background, campus activities, and personal experiences and attitudes can be used to build a predictive classifier to identify students at risk of binge drinking. By dividing the survey data into five categories of data, organized by how difficult they are to collect, we also explore how well our classifiers perform using limited information.

We begin by introducing the dataset and our methodologies for processing the data and evaluating predictive classifiers. From there we move on to detailed explorations of the techniques we have used, including Generalized Linear Models, clustering techniques like k-means, and PCA. Finally, we use the insights gained from using these techniques to provide an overall evaluation of our

work and how it impacts the problem of binge drinking on college campuses.

2 Dataset

Our analysis is based on the data from the 2001 Harvard School of Public Health College Alcohol Study [?]. This study surveyed 10,904 students from 119 different four-year universities about their drinking patterns, attitudes towards alcohol and other drugs, and participation in various aspects of campus life, among other things. The features used in our analysis come from the pool of 342 survey questions, ranging from basic demographic information ("How old are you?") to detailed information about a student's perception of their own campus ("If a student is caught on your campus using a fake ID to get alcohol, what is likely to happen to the student?").

We are interested in predicting whether a student is likely to engage in binge drinking. In accordance with the survey, we will define binge drinking as consuming five or more drinks in a row. We thus wish to predict the student's response to survey question C1: "Think back over the last two weeks. How many times have you had five or more drinks in a row?", with possible responses "none", "once", "twice", "3 to 5 times", "6 to 9 times", and "10+ times". The other survey responses will serve as features for our classifier.

3 Methodology

FILL ME IN!

4 Generalized Linear Models

FILL ME IN!

5 Clustering

FILL ME IN!

6 PCA, ICA, SVD

What is these even?

7 Final Evaluation

FILL ME IN!

8 Conclusions

Conclude meh!